Intel Corporation

Investment memo · NASDAQ: INTC · 24 October 2026

Produced by **Bryant Tan**

ONE-MINUTE VIEW

Intel's reversal story is playing out. It slipped at 10nm, when its manufacturing process failed and took its products with it. It has since bet hard on coming back strong: it took High NA, the newest chip-printing tool, before anyone else, committed to 14A, and is raising capex to build it. 18A, its current process, is the proof, in volume since the end of 2025. 14A, the next one, is when Intel competes on a level field with the winners of this period, arriving in 2028, the same year as TSMC's A14.

The market re-rated Intel on a narrative, not on the technology. Intel only got re-rated after the government invested and the CPU shortage narrative took hold. Most analysts still haven't warmed to an unsexy stock: two-thirds rate it a hold, and they're waiting on gross margin and outside customers. So the market prices 14A as an option that pays only once an outside customer signs.

I think they misunderstand how strong Intel is again at the leading edge. It's level with TSMC on the transistor. It's ahead on backside power, which feeds power through the back of the chip, and on High NA and in packaging. Even TSMC has adopted the idea Intel shipped first: since 2024 it has built small silicon bridges into its packages to join chips together, which Intel has done since 2017. I see 14A as the node that renews Intel's premium, the higher price its newest chips earn, and Intel's own products pay for it first. I don't need the shortage or an outside customer for this to work. An outside customer at scale is the Bull.

Intel can be paid whoever designs the chip. It's an IDM, a chipmaker that designs and makes its own chips, and it also makes other companies' designs, as TSMC does. So it can design, make and package a custom chip and keep every layer of its revenue, or sell just the wafer and the package to a customer that brings its own design. Very few companies can do all three, and that lets Intel capture more than one market.

Why now. 2027 is the test, on 18A's revenue and margins. I expect 14A's first design win by the end of 2027: a custom chip Intel designs for a customer, once the design kit, the toolkit chip designers work with, is in customers' hands. 2028, with 14A, is the turnaround.

What would prove me wrong. Intel not being able to sell at a premium: 14A slipping, so 18A ages with nothing above it as the shortage ends, and Intel's prices falling once supply catches up. Share on its own isn't the test. If AMD takes share but Intel sells fewer chips at a higher price, that isn't necessarily bad, and an Arm chip that Intel designs and makes is still Intel's revenue.

BUY

Base $133.23 · 53.2% above the price

PRICE AT RATING

$86.96

the close on 23 October 2026

VALUE PER SHARE

Bull	$212.10	+143.9%
Base	$133.23	+53.2%
Bear	$47.54	−45.3%

The DCF by part ending at the competitors' multiples

HORIZON

About three years, to 2029

The first year 14A's products sell

MARKET DATA AT THE PRICE

Market value	**$481.0 billion**
Enterprise value	**$479.2 billion**
Diluted shares	**5,531 million**

From the valuation model, at the close of 23 October 2026.

Variant view

What the market believes. The market has re-rated Intel on the government's investment, the CPU shortage and the hope of a foundry. The stock is up 136% since the start of 2026, and at about 51 times the earnings analysts expect, it pays now for earnings Intel has yet to report. The analysts haven't warmed to it: two-thirds rate it a hold. The street reads 2026 as a cyclical gift: CPUs are short, prices are up, and AMD and Arm share the gain. It wants gross margin on a path to 50%, treats 14A as an option that pays only once outside customers sign, and won't pay for Foundry until they do.

What the market fears now. In October the market turned wary of anything tied to AI, and the fear is very real. I'm not saying to ignore it. Far more moves the market than this memo can explain, but three reasons come first:

- **Money costs more.** The war with Iran and the Houthis' attacks on Saudi Arabia have pushed oil up, Brent to almost $117 a barrel on 23 October, its highest since April, and with it inflation and the rate rises markets expect. Government bond yields reached their highest in years in early October: the US ten-year above 5%, Germany's the highest since 2008.
- **AI's revenue is unproven.** OpenAI told investors in October that its annualised revenue was almost $50 billion, about $20 billion below the figure reported weeks before. Gartner reckons AI data centres need $8.5 trillion of revenue by 2030 to pay for what's announced and earn a return, against no more than $3.6 trillion it expects, and sees more than 80% of the neoclouds, the AI cloud companies, ceasing to operate on their own.
- **The build-out is borrowed.** In mid-October, insuring $10 million of CoreWeave's debt for five years cost more than $1.1 million a year, a price that implies about a two-in-three chance of default, and Oracle's about $300,000. Even Amazon's, Alphabet's and Meta's protection traded at record prices.

The market's timeline.

- **October 2026:** 14A's 0.9 design kit, the one outside customers need to design chips for 14A, due to reach them.
- **2027:** capex up again; customers ramping on EMIB-T, Intel's packaging for the largest AI chips; 14A's risk production, its first runs before volume, in the second half.
- **End of 2027:** Foundry at breakeven, Intel's target. I expect it in 2028, as 18A matures.
- **2028:** 14A in high volume, alongside TSMC's A14.

Where I differ.

- **On the technology (Pillar 1).** The market prices 14A as an option. I price it as the renewal of Intel's premium. A leading-edge node is the flagship for about three years, then earns for five to ten more as a mature node. 18A's premium runs until 14A arrives: its wafers in 2028 and its products from 2029, paid for first by Intel's own products. Intel's lead in High NA, backside power and its own packaging makes 14A the node where Intel stops catching up and starts setting the pace.
- **On the place (Pillar 2).** The market reads Washington's backing as a moat in principle, not as revenue. I think it turns into wafer and packaging orders before TSMC's and Samsung's US fabs can offer the same, and Apple's wafers would be the first.
- **On demand and the business model (Pillar 3).** The market agrees that AI agents drive CPU demand, but prices it as a tide that lifts AMD and Arm as much as Intel. I think Intel takes more of it than the market credits. It still ships about seven of ten x86 processors and two of three x86 server processors. Companies have run on software built for x86 for decades. And Intel is one of very few that can design, make and package a custom chip and keep every layer of its revenue. A cloud that wants its own Arm chip can build it with Intel, too.
- **On AI itself.** I acknowledge those risks; some of the build-out will lose money. But the deeper question is how you price something with recursive self-improvement, AI that helps build better AI. The railways and electricity had nothing like it: once the track or the grid was built, what it could carry was known, and that's

the arithmetic the market applies now, revenue against capex to 2030. AI's reach compounds, into the physical world as much as the digital, from factories and labour to software and security. No model prices that, and mine doesn't try. So I'd rather own the layer that's paid however it plays out: Intel designs and makes the chips, a pick and shovel with a turnaround's upside. And when customers design their own chips, as OpenAI now does (Pillar 3), it's the merchant designers' margins at risk; the fab is paid either way.

What would prove the market right: no outside 14A customer named through 2027; Intel's prices falling as supply catches up, or as the clouds cut their AI spending, in 2027, and again in 2028 with no 14A product above them; Washington's support staying rhetoric. AMD and Arm taking share proves it right only if Intel's prices fall with it.

The case

The loop and its clock. Intel's dominance was a loop: a process lead made better chips, better chips won share, share filled the fabs, and full fabs paid for the next process. At 10nm the loop ran in reverse: the process failed, the products tied to it stalled, and AMD took share by building at TSMC. The loop also runs on a clock. A leading-edge node is the flagship for about three years, carrying the highest prices. It then lives on for five to ten more as a mature node, cheaper but highly profitable, making the chips that don't need the newest process. The flagship takes the price, the mature node takes the margin, and the mature margin pays for the next flagship. That's how TSMC runs, and it's what Intel lost at 10nm: the flagship arrived years late, so nothing renewed the premium and nothing paid for the node after it.

I think the loop is turning back. 2027 is the test, on 18A's revenue and margins; 2028, with 14A, is the turnaround. By 2030 it's in Intel's reported numbers: share held after supply has caught up, custom silicon well beyond its 2026 size, and Foundry's fabs full of Intel's own 18A and 14A wafers, past breakeven and converging on TSMC's margin. 18A is the flagship to 2028. 14A takes over in 2028, its wafers first and its products from 2029, and holds the premium to 2031, with 10A and 7A in development behind it.

Two premiums. The CPU shortage of 2026 is a tailwind for Intel and for the whole industry. I take it, but I don't rest on it. Two premiums sit in Intel's 2026 prices. The larger is the product's: chips on the new nodes replacing those on Intel 7, its older process, at higher prices, most of the rise on Intel's own account. It rises as the technology advances, most in each node's premium years. The smaller is the shortage's: demand pricing on every node, taken partly to cover dearer substrates and memory, which lasts until supply catches up at the end of 2027. My case rests on the first, and 14A renews it. A third premium is Foundry's, the price of a wafer, which only outside volume earns. That one belongs to the Bull.

Base: a re-integrated Intel. Intel's own products fill 18A and 14A, and that is customer demand. It means processors sold through the PC and server makers and to the clouds, short of supply. It also means custom chips that DCAI, Intel's data-centre business, designs for a customer and makes or packages. Intel says its committed demand, from its own product roadmap and from outside design wins, will set the scale and pace of 14A's fabs; this case sizes them to its own roadmap. The first 14A design win comes by the end of 2027 and books in DCAI. In 2028 14A's wafers enter volume and 18A-P, the tuned version of 18A, holds the top of Intel's range while 18A moves down into the mainstream; from 2029 the top is 14A's. The first outside chips on 18A-P ship in 2027 and reach volume in 2028, Apple's lowest-end Mac and iPad chips among them. So Intel's price and mix hold through the handover, and the 18A lines stay full. Foundry breaks even on yield and volume in 2028 as 18A matures, a little after Intel's target. Its margin converges toward TSMC's: under way by 2030, and settled in the years after.

Bull: the merchant foundry. The Base, plus outside customers on 14A and on EMIB-T packaging, taking Foundry to the scale a leading-edge node needs by 2030. Foundry takes TSMC's shape: a flagship node for the top customers and a mature node for the long tail. Its margin is close to TSMC's by 2030, on the volume the outside customers bring. A customer that brings its own design to 14A, Apple among them, is Foundry's, and belongs here.

Bear: the muddle-through. Intel completes 14A, as it committed to in Q2 2026, and fills it with its own top tier and nothing more. No outside design win comes, for DCAI's custom chips or for Foundry, so the expansion runs at the pace Intel's own roadmap can justify. Meanwhile supply catches up, the demand pricing of 2026 goes, and the depreciation of fabs built for more arrives. 18A ages into a mature node without the volume a mature node needs. Foundry's breakeven slips to 2029 or 2030, and the premium is renewed on too few wafers to pay for the node after 14A. The tail of this case is the fallback: if the fill falls far enough, Intel pauses 14A, a step its 2025 10-K warned may be effectively irreversible, buys from TSMC and runs Products as a fabless company, settled by 2030.

What the case requires.

- **Spending ahead of the earnings.** Intel expects more than $20 billion of capex in 2026 and significantly more in 2027, while 14A's products only arrive in 2029. I see that as the price of the model: without the next node there's no premium to renew, which is what 10nm cost Intel. It's safe because Intel is its own biggest customer and paces 14A to committed demand, its own products first. It turns into the Bear if Intel keeps raising capex while its own premium chips sell less: fabs built for more, depreciating on too few wafers.
- **The owners pay for it: a ticket, not a break.** Nothing forces a raise. June's $29.7 billion of cash and short-term investments plus the August sale roughly match Intel's $50.5 billion of debt, and it still sold $6.5 billion of new bonds in the second quarter. It chose to sell shares while they were dear, which costs its owners less than the partner money it took at its low point: in April it paid $14.2 billion to take back the half of Fab 34's profit it sold to Apollo in 2024, a buy that only pays if its own products keep the fab full. The share count is up about a fifth since the end of 2023, and the August sale, the government's escrowed shares and employees' unvested units add about another tenth, so every case I value carries the dilution.
- **The team has to keep delivering:** 18A at volume, 14A on schedule, EMIB-T into production. No figure covers that, and everything else depends on it. What I need from management is capital spent only against committed demand, and the nodes on time. On the first, Lip-Bu Tan has broken Intel's habit of building ahead of demand: he called that spending "unwise and excessive", cancelled Germany and Poland, slowed Ohio, and adds capacity only against volume commitments. His biggest call is my case: in July 2025 Intel warned that 14A might not pay without an outside customer, and in 2026 it committed to 14A on its own products' demand. 2027's capex is where the discipline is tested. The second is half-proven: 18A reached volume on time, EMIB-T slipped into 2027, and 14A is the real test. He started with the culture, halving the management layers, which is why I watch turnover through the ramp. He's paid in options and performance shares, with $25 million of his own in the stock. The risk is how much rests on him: in August 2025 the President called for him to resign after a senator questioned his ties to China.

The true tests are Foundry's cost per wafer and Intel's gross margin. Foundry's own profit moves with the price Intel charges itself, so it isn't a test on its own.

Pillar 1. The Irony Reversal – Technology & Process Gains, 14A is the Inflection Point

Intel slipped because it asked old tools for a new node. At 10nm, EUV, the tool it wanted for the finest layers, wasn't ready when it had to lock the node's design. So it set a density goal far beyond the industry's on the older tools, and the defects from printing each layer several times stalled the node for years. Now the roles have reversed. TSMC is building its next nodes without High NA, the newest chip-printing tool, and keeping it out of volume until about 2030 on its cost and maturity. Intel took it first. Intel lost 10nm stretching the old tools before the new one was ready; the leader is now stretching them by choice, with the new one available.

That's why I think the market has 14A wrong. It isn't an option on Foundry. It's the node where Intel's leads come together: High NA, the second generations of its transistor and of backside power, and its own packaging. It arrives in 2028, the same year as TSMC's A14 and within 5% of it on performance by Intel's own account. It's where Intel stops catching up and starts setting the pace, and that's what renews the premium.

Intel has the depth for it. It has set the industry's turns rather than followed them: the first FinFET, the silicon bridge embedded in the package, the first node to pair gate-all-around transistors with backside power, High NA in production. Behind that record is the research budget of a foundry and a chip designer in one: $79 billion over the five years to 2025, more than TSMC and AMD spent combined, at $58 billion, and over two and a half times TSMC's alone.

- **The tool.** Intel is the only chipmaker using High NA in volume, on 18A layers since 2026, about two years before TSMC or Samsung plan it for logic. High NA prints in one pass what older tools need several for, so each wafer takes fewer steps and has fewer chances for a defect. Intel is learning it in production on a proven node first, the way TSMC brought EUV in.
- **The wafer.** RibbonFET, Intel's gate-all-around transistor, wraps the gate around four stacked silicon ribbons on all four sides, where a FinFET covers three. So it leaks less and drives more current in the same space. On shipping clocks, 18A matches or beats TSMC's N2 on speed, while trailing on density. On top of it sits backside power, which feeds power through the back of the chip. 18A is the only node in volume pairing the two until 2027, TSMC's main node for 2028 has no backside power at all, and 14A brings the second generation of both.
- **The package.** Intel pioneered the embedded bridge: patented in 2008, shipping since 2017, and joining 17 tiles with 12 bridges in Clearwater Forest. It puts silicon only where two dies meet, where TSMC's CoWoS lays a slab under the whole package. That costs a few hundred dollars a chip against $900 to $1,000, on Bernstein's estimate. TSMC followed seven years later with bridges of its own, and is reported to be developing a package it calls "EMIB-like" internally. SemiAnalysis, which still rates CoWoS ahead on several counts, calls EMIB-T the most credible alternative. The next layer is already in Intel's labs. It started the race to glass cores, the stiffer base the largest packages will need, with more than a decade of research, and has embedded its bridges in glass. No product uses glass yet, and Intel points to about 2030, but it shows how deep the research runs.

So far TSMC is making its choice work, on a smaller leap and on schedule. That its extra steps become a handicap at the nodes after A14 is my view, not yet the evidence. What the evidence shows is Intel back at parity on the transistor, ahead on the tool and on backside power, and running ahead of plan: 18A's yields since April 2026, and 14A's defects falling faster than on any Intel node since 22nm.

Status, 24 October 2026. On track on the wafer, later than planned on the package. 18A is in volume with its yields rising, 18A-P reached risk production on schedule, and EMIB-T's customer ramps moved from late 2026 to 2027. The 0.9 design kit is due to reach outside customers this month, and Intel hasn't yet said it has.

Next test. 14A's design kit and risk production on time, by the second half of 2027, and 18A's yields at industry level by its end. A named outside 14A customer tests the Bull, not this pillar.

Pillar 2. Intel's Culture Learnings & Geopolitical Advantage

The AI race runs between America and China, and America's side of it runs through Taiwan. Almost every leading AI chip is made by TSMC, whose customers are mostly American and whose plants are mostly in Taiwan. That leaves TSMC caught between the two powers: Washington pulls its production onto US soil with tariffs, while Beijing rehearses blockades around the island. Intel carries far less of this. It's the only company doing both leading-edge chip research and high-volume manufacturing in the US, and two-thirds of its plant is there.

I think that turns into orders, not just goodwill. The market reads Washington's backing as a moat in principle, not as revenue. I think it becomes wafer and packaging commitments before TSMC's and Samsung's US fabs can offer the same, and the window is open now.

- **Policy is built around Intel.** The government became a shareholder in 2025, with warrants that make it costly for Intel to give up control of Foundry. It also struck the milestones from Intel's CHIPS agreement while TSMC's remain. Its trusted route for defence chips, Secure Enclave, is designed around Intel, and the tariff design and a 35% investment credit reward whoever already makes chips at home. The ties are

forming: NVIDIA and SoftBank put in $5 billion and $2 billion, and after the Commerce Secretary pressed Apple, the President announced it would work with Intel.

- **Apple is the test.** Apple designs its own chips and makes none, so an Apple deal means Foundry wafers, not a chip sold through DCAI. The analysts who follow its supply chain put its lowest-end Mac and iPad chips on 18A-P, shipping from mid-2027 and in volume in 2028. That's second-source wafers at the Base's small scale. Some non-Pro iPhone chips on 14A from 2028 would be the Bull. Apple is also the customer most able to wait, as the first and biggest customer of TSMC's Arizona fabs.

Culture is the part the market misses. People treat companies as if they were machines, but every business problem is a human problem, and a culture that doesn't fit the place it works in is destructive. Chipmaking shows it clearly. TSMC's edge rests on its people. Its founder credits Taiwan's work culture, with fab staff who stay on call, for keeping Taiwan ahead, and TSMC can reassign thousands of engineers between its sites when it needs to. That model is TSMC's strength at home and its problem in America.

- **The mismatch is already costing Intel's rivals.** TSMC delayed its first Arizona fab in 2023 for want of skilled workers, and about half of its 2,200 Arizona staff in 2024 had been sent from Taiwan. US applicants and employees have sued it alleging discrimination, though the claims are unproven. Korean manufacturers rotate engineers through short visas, and the 2025 immigration raid on a battery plant in Georgia detained 475 people, more than 300 of them Korean, and threatened to delay Samsung's Taylor fab. The President's answer to foreign investors was that they "hire and train American Workers".
- **The rivals can pay to close the gap, and they are paying.** TSMC's chairman said in January that Arizona's yield and defect density were "almost equal to Taiwan", after "we worked hard". It costs: TSMC expects its overseas fabs to cut its gross margin by 2 to 3 points in the early years and 3 to 4 later. US labour cost about 40% more than Taiwan's on a 2020 estimate, and the US could be short of up to 157,000 chip engineers and technicians by 2030.
- **Intel has no culture to transplant.** It has run fabs in America since 1976, in Oregon and for more than 40 years in New Mexico, built around an American workforce that costs more and has choices. Its bet on tools over hands fits that workforce: High NA prints in one pass what older tools need several for, so each wafer takes fewer steps. That this means fewer people per wafer is my reading; no one publishes engineers per fab.

The same lens applies to Intel. If culture is paramount, Intel has to pass its own test. Its undesired turnover rose to 7.9% in 2025 from 5.9% in 2024, after its layoffs, and Intel itself warns that restructuring can hurt morale and cost it expertise. It ties part of its bonuses to morale, engagement and retention. Keeping its engineers through the 14A ramp is the culture test I'd watch.

Two limits keep the place argument honest. Once the tariff wall is built, it rewards any US fab, so TSMC and Samsung earn the same relief as their US fabs ramp. And Intel's customers are mostly abroad, with a quarter of its billings in China. So the place pays only once the technology of Pillar 1 is there: US demand follows the best process, not the address.

Status, 24 October 2026. Holding, on a clock. The stake is in place, broader tariffs are promised but not proclaimed, TSMC's US build keeps growing, and no customer has confirmed a government-brokered deal.

Next test. A customer confirming a government-brokered deal, Apple first, with its 18A-P chips shipping from mid-2027. The window closes for Apple when TSMC's Arizona 2nm fab reaches volume. And Intel's undesired turnover coming back down as 14A ramps.

Pillar 3. The Full Stack – Design, Manufacturing & Packaging, Agentic AI Is the Demand Engine

AI's money is moving from training models to using them. Training needs thousands of GPUs, the chips that train and run the models. Using the models is what gets paid for, and more and more that means agents: AI software that takes a task and works it through on its own, calling a model many times along the way. Each call runs on a GPU or a similar AI chip. The work around the calls runs on CPUs: the tools the agent uses, the

code it runs, the data it handles and the steps it coordinates. And more of the work is moving to chips built for one job.

I think the market has the demand right and Intel's share of it wrong. It agrees agents drive CPU demand, then counts the chips and splits them between Intel, AMD and Arm. I count who gets paid. Every one of those chips has to be designed, made and packaged, and Intel is one of very few companies that can do all three. So it gets paid whoever designs the chip: for the finished chip where Intel designs it, for the wafer and the package where the customer does. That's why share on its own isn't my test.

- **The demand runs on CPUs, and Intel can't make enough.** Intel puts agent work at one CPU per GPU or better, against about one per eight GPUs for training. That's a seller's number, but the buying bears it out. Demand for Intel's chips ran above what it could supply through the first half of 2026, and in September its chief executive said it could supply only half its customers. Some of that buying is stock-building ahead of the shortage, on IDC's reading. That's why 2027, as supply catches up, is the test.
- **The enterprise runs on x86, and Intel follows the clouds to Arm.** x86 is the chip standard Intel's and AMD's processors share, and companies have built their software on it for decades. Moving off it costs a company in ways a cloud that writes its own software doesn't face. Intel sells to companies through the PC and server makers, and its efficient server chips compete on performance per watt, built on what AMD doesn't yet have: the backside power and bridges of Pillar 1. The clouds are the harder customer, because they design their own processors on Arm, a rival chip standard they license. Intel designs those too. The infrastructure processor it built with Google, which takes a server's networking and storage off its main processor, runs on sixteen Arm cores and sells through DCAI at the price of the finished chip.
- **Designs multiply, fabs don't.** The clouds and AI companies now design chips for the work they run, and fast: OpenAI's first, built with Broadcom, went from design to tape-out, the final design sent to the fab, in nine months on their own account. Making them isn't spreading. TSMC still makes almost all of them, and its packaging is so short that its chairman says it limits his customers' growth. So every new design joins the same queue. The wafer and the package are the scarce layers, and Intel sells both to any designer, even packaging chips made elsewhere. Where Intel designs the chip too, as it will for Fortinet's next security processor, it earns every layer. Its custom-chip business is small, but it nearly tripled in the year to the second quarter of 2026. With customers learning to design for themselves, the design's real value to Intel is the wafers it pulls into Intel's fabs: the old loop, extended to other companies' chips.

Two limits keep the stack argument honest. What the scarcity brings Intel is volume, not the designer's margin: design keeps its premium, and NVIDIA still earns a higher gross margin than TSMC. And Foundry is a division Intel owns outright, so a customer that competes with Intel's products has to trust a rival with its designs. Intel names that concern itself, and TSMC's neutrality is built to avoid it. No customer is reported to have turned Intel down for it, and NVIDIA has invested in Intel, but with few outside customers so far, it's still open.

Status, 24 October 2026. Strengthened, durability untested. CPU demand runs ahead of supply, and custom chips are growing fast from a small base. The test comes in 2027, as supply catches up.

Next test. Intel's prices holding as supply catches up through 2027: the premium, not share on its own. A chip Intel designs, makes and packages for an outside customer, named and in volume. And independent tests of Clearwater Forest, Intel's newest efficient server chip, which Intel's own tests put ahead of AMD's densest on performance per thread per watt.

The value and the call

The call: Buy, with a horizon of about three years, to 2029. My preferred method, a DCF of Intel's two parts that ends at their competitors' multiples, puts Intel's value against its $86.96 close of 23 October 2026 at:

- **Base:** $133.23, 53.2% above the price.
- **Bull:** $212.10, 143.9% above.
- **Bear:** $47.54, 45.3% below.

The market doesn't pay for Foundry. At $86.96 it pays about $48 billion less for all of Intel than Products alone is worth at its competitors' multiple of 2027's profit. The October selloff moved the price, not my value: my multiples are medians through a cycle that takes in 2022's bust. Three years is how long my case takes to prove: 2027 tests 18A and the premium, 2028 brings 14A's wafers and Foundry's breakeven, and 2029 brings 14A's products, when the case is judged.

How I value Intel. I value Intel's two parts, Products and Foundry, on the model's forecast to 2030, each at its own cost of capital: 13.5% for Products, on the risk of designers like AMD and Qualcomm, and 11.2% for Foundry, on the risk of TSMC and GlobalFoundries. The rest of the group keeps its cash-flow value. The total reaches a value per share through Intel's debt, cash and other claims, on a share count that includes the dilution. Four methods run on that forecast, and they differ only in what they assume:

Method	What it assumes	Bear	Base	Bull
DCF by part, ending at the competitors' multiples: preferred	Intel's cash to 2030, then each part valued like its competitors	47.54	133.23	212.10
Sum of the parts	Each part valued like its competitors today, on 2027's results	87.82	111.44	121.85
DCF by part, steady growth	Intel's cash to 2030, then no return above its cost of capital, forever	6.54	44.03	81.63
DCF of Intel as one business	The same, at one cost of capital for the whole company: the check	8.33	42.84	76.88

Why this method. An intrinsic DCF needs a business at steady state when its forecast ends, and Intel won't be there by 2030. Foundry will still be climbing to its margins, and Intel's cost of capital will still be a turnaround's. The steady-growth DCF capitalises 2030 as if it were steady: it keeps today's cost of capital forever, 13.5% for Products, and assumes Intel earns no more than that on what it invests. That values Products at about six times its profit, where its competitors trade near 24, and it accounts for $69 of the $89 between that DCF and mine.

I could stretch the forecast past 2030, but my case goes no further than 14A. So I end where I can see, and value each part at what the market already pays businesses that have finished the climb. That puts my value $22 above the sum of the parts, which prices 2027's results today. Mine adds the cash Products earns to 2030, and it values Foundry on its 2031 earnings, after its losses on the way, where the sum of the parts can only value it on revenue.

The comps. Each multiple is the median of what two competitors were worth at each year-end from 2021 to 2024, against their next year's results, a span that takes in a boom, a bust, a recovery and the AI run:

- **Products, on operating profit:** AMD (34.7 times) and Qualcomm (13.4), a median of 23.8. Arm, at 151 on its one listed year-end, is shown but not counted.
- **Foundry, on EBITDA once it earns:** TSMC (8.7 times) and GlobalFoundries (11.4), a median of 10.4, with UMC beside.
- **Foundry, on revenue while it loses money:** GlobalFoundries (4.1 times) and UMC (2.3), a median of 2.97, with TSMC at 6.0 as the ceiling.

What the value rests on. About 92% of the value comes from the 2030 sale, so the case has to earn the multiples. Three things move it:

- **What 23.8 times asks of Products.** It's the median of two very different businesses. At Qualcomm's 13.4 the Base would be $91.60, still above the price; at AMD's 34.7, $176.92. To be worth 23.8 times at a 13.5% cost of capital, Products' profit has to grow at least 10% a year after 2030, and more if that growth takes investment. The model grows it about 19% a year from 2027 to 2030. The case has to keep it going: each node renewing the premium, 14A to 2031, with 10A and 7A behind it.
- **Foundry's path.** The method sells Foundry on its 2031 earnings, at 10.4 times EBITDA, because by then it earns. If it were still losing money, the sale would be on revenue at 2.97 times, and the Base about $11 lower. That's why its breakeven in 2028 matters.
- **The price Intel charges itself for wafers.** Products' profit carries 23.8 times and Foundry's 10.4, so every dollar of the bill moved from Products to Foundry lowers the total. I set it where market wafer prices put it.

Catalysts

These are the events I expect to bring the market round to my view, in the order they come. Most fall in 2027, the year the case is tested, and the one I'd weigh most is Intel's prices holding as supply catches up.

When	Catalyst	What it proves
Mid-2027	Apple officially announces its 18A-P chips with Intel	Foundry has a marquee outside customer, which the market won't pay for until it sees one
2027	Customers ramp on EMIB-T	Intel's packaging wins AI chips that TSMC can't package fast enough
Through 2027	Intel's prices hold as supply catches up	The premium is the product's, not the shortage's
Second half of 2027	14A's 1.0 design kit and risk production on time	The node that renews the premium is on schedule
By the end of 2027	14A's first design win, a custom chip Intel designs for a customer	14A is a node customers design for, not an option
By the end of 2027	18A's yields at industry level	Gross margin can climb toward the 50% the market wants
2028	Foundry at breakeven, and 14A's wafers in volume alongside TSMC's A14	Foundry stops costing the group money, and Intel competes on a level field
2029	14A's products at the top of Intel's range	The premium is renewed, and the case is judged

An outside customer at scale on 14A would be the Bull's catalyst, not the Base's.

What would prove me wrong

I rank these by what they'd do to my case. Where a break needs a model figure, it's the Bear's path: if Intel's numbers track the Bear's, I'm wrong.

	What would prove me wrong	Period	Breaks if
1	**14A misses its dates**, so the premium isn't renewed	Q4 2026	The 0.9 design kit hasn't reached outside customers
		January 2027	The FY2026 10-K backs off the commitment to complete 14A
		H2 2027	The 1.0 design kit or risk production slips into 2028
		2028 to 2029	14A's wafers miss volume in 2028, or its products miss the top of Intel's range in 2029, so 18A ages with nothing above it
2	**Intel can't hold its premium** through the handover	2027	As supply catches up, or the clouds slow their AI spending, Intel's prices stop rising: the premium tier up no more than the 4% carried over from 2026, the Bear's path. First sign: finished goods back above 2024's 36 days of cost of sales, or a reserve on the premium parts
		2028	The premium tier's price falls with no 14A product above it, tracking the Bear's 10% fall
3	**The true tests track the Bear:** Foundry's cost per wafer and Intel's gross margin	Each quarter of 2027	The drag from 18A's dearer wafers on Foundry's profit stops shrinking, from $489 million in the first quarter of 2026 and $340 million in the second
		By the end of 2027	18A's yields still short of industry level
		By the end of 2028	Foundry still short of breakeven, judged on its cost per wafer, not the price Intel charges itself
		FY2027 and FY2028	Gross margin on the Bear's path, about 33% in both years
4	**Capex runs ahead of committed demand**	2027	Capex keeps rising while customer deposits, long-term agreements and designs on 14A don't, or while Intel's premium chips sell less
5	**No one outside signs**	By the end of 2027	No 14A design win, for a DCAI custom chip or for Foundry
		By the end of 2027	No outside chip shipping on 18A-P, Apple's lowest-end Mac and iPad chips first
6	**Intel loses its engineers** through the ramp	January 2027 and 2028	Undesired turnover rises again from 2025's 7.9%

Three risks have no date and could break the call at any time: the Israeli fab, which makes Intel 7 and is uninsured against war; China, a quarter of Intel's billings, restricting Intel further; and Washington's support staying rhetoric.

If you'd like to discuss any of this further, reach out to me.

Disclosure. This memo is personal research. It is not investment advice, or an offer or solicitation to buy or sell any security. As of 24 October 2026, I hold a position in Intel.